

# First-Principles AI Study Plan — Rust Edition (20-40h)

## Style

Specs & papers only · slow thinking · implementation-first · no frameworks

## Goal

Rebuild the primitive stack behind modern AI in Rust the same way you rebuilt ALUs, hashes, and VMs.

Focus on **execution, memory, invariants** — not performance, hype, or tooling.

---

## Scope

You will rebuild, from scratch:

- Math primitives → gradients → backpropagation
- Attention → multi-head attention
- One Transformer block (forward pass)

Optional extensions add efficiency and control layers.

---

## Final Outputs (non-negotiable)

By the end you must have:

1. A pure-Rust MLP trained via **manual backprop** (XOR + small classification)
2. A pure-Rust attention implementation (Q/K/V + masking + stable softmax)
3. A Transformer block **forward pass** (attention + FFN + residuals + layer norm)
4. Gradient-check tests (finite differences) for MLP and attention

If it's not testable, it's not understood.

---

## Tooling Rules

- Language: **Rust only**
  - Data structures: `Vec<f32>` / `Vec<f64>`
  - Optional helper crate: `ndarray` (arrays only, **no autograd**)
  - No ML frameworks
  - No hidden gradients
  - No GPU
-

# Phase 0 — Math Foundations (4h)

## Concept

Linear algebra is executable structure.

## Read (slowly)

- Vectors, matrices, dot products
- Gradients, Jacobians, chain rule
- Log-likelihood, cross-entropy intuition

## Supplement (read only if stuck)

### • Book

*Deep Learning* — Ian Goodfellow

- Chapter 2: Linear Algebra
- Chapter 6.2–6.5: Gradients, Jacobians, chain rule

### • Blog (text)

Andrej Karpathy — *Micrograd: The spelled-out intro to neural nets*

(Focus on math + computational graph explanation, ignore Python details)

## Build (Rust)

- Matrix multiplication (nested loops)
- Stable softmax
- Cross-entropy loss
- Finite-difference gradient checker

## Deliverable

```
math_primitives/
├── matmul.rs
├── softmax.rs
├── loss.rs
├── gradcheck.rs
└── tests/
```

# Phase 1 — Backpropagation (6h)

## Concept

Backprop is the learning algorithm. Everything stacks on it.

## Read

- *Learning Representations by Back-Propagating Errors*

Rumelhart et al., 1986

(Focus on algorithm flow, not history)

## Supplement

- **Book**  
*Neural Networks and Deep Learning* — Michael Nielsen
- Chapter 2: How the backpropagation algorithm works
- **Blog (text)**  
Andrej Karpathy — *Yes, you should understand backprop*

## Build (Rust)

- 2–3 layer MLP
- Activations: sigmoid + ReLU
- Loss: MSE → cross-entropy
- Manual backward pass using chain rule
- Gradient checking on small networks

## Deliverable

```
mlp_from_scratch/
├── forward.rs
├── backward.rs
├── train.rs
└── gradcheck.rs
└── tests/
```

---

# Phase 2 — Attention (6h)

## Concept

Attention replaces recurrence with pure linear algebra.

## Read

- Attention chapter from a deep learning textbook  
(Focus on Q/K/V, scaling, masking, multi-head logic)

## Supplement

- **Book**  
*Dive Into Deep Learning*
- Chapter 10.1–10.3: Attention mechanisms
- **Blog (text)**  
Jay Alammar — *The Illustrated Transformer*

## Build (Rust)

- Scaled dot-product attention
- Causal masking
- Stable softmax

- Multi-head attention  
(reshape → attention → concat → projection)

## Deliverable

```
attention_from_scratch/
└── attention.rs
└── multi_head.rs
└── mask.rs
└── gradcheck.rs
└── tests/
```

# Phase 3 — Transformer Block (4h)

## Concept

Attention + MLP + normalization + residuals

## Read

- *Attention Is All You Need*  
(Architecture sections only)

## Supplement

- **Book**  
*Deep Learning* — Ian Goodfellow
- Chapter 9.3: Normalization
- **Blog (text)**  
Sebastian Raschka — *Layer Normalization Explained*

## Build (Rust — forward pass only)

- Token embeddings
- Positional encoding (sinusoidal)
- Multi-head attention
- Residual connection + layer norm
- Feed-forward network
- Residual connection + layer norm

## Deliverable

```
transformer_block/
└── embeddings.rs
└── positional_encoding.rs
└── transformer_block.rs
└── tests/
```

# Extended Track (Optional — up to 40h total)

Only start after completing core phases.

---

## Extension A — Train a Tiny Transformer (8-10h)

### Tasks

- Toy tasks (copy, next-token prediction)
- Rust only
- Expect pain

### Supplement

- **Book**  
*Deep Learning* — Ian Goodfellow
  - Chapter 8: Optimization for Training Deep Models
  - **Blog (text)**  
Andrej Karpathy — *A Recipe for Training Neural Networks*
- 

## Extension B — Distillation (4-6h)

### Tasks

- Teacher → student setup
- Soft targets with temperature
- KL loss + supervised loss

### Supplement

- **Blog (text)**  
Hugging Face — *DistilBERT Explained*
  - **Book**  
*Deep Learning* — Goodfellow
  - Chapter 7.3: Regularization as information control
- 

## Extension C — Control / Alignment (4-6h)

### Tasks

- Reward model (pairwise ranking)
- Simple policy improvement loop

## Supplement

- **Blog (text)**  
Lilian Weng — *RLHF Explained*
  - **Blog (text)**  
OpenAI — *InstructGPT* (methodology sections only)
- 

# Core Papers & How to Use Them

### Rule

- If a paper defines a primitive → **build it**
- If it defines a policy or control layer → **understand it and move on**

### Required

- *Learning Representations by Back-Propagating Errors* (1986)  
→ Derive equations, then code backprop
- *Attention Is All You Need* (2017)  
→ Treat as a CPU architecture spec

### Optional

- *DistilBERT* (2019) — optimization layer
  - *RLHF* (2022) — context only, read-only
- 

# Success Criterion

You can **re-derive modern AI from math and Rust code**, not belief, tooling, or APIs.

If you can build attention with explicit loops, **you understand the system**.